

L A N D O N O S

Performance Research Cockpit

User Guide

An AI-guided research training cockpit that teaches and structures responsible compliance and business research. A luxury instrument cluster reads your work at a glance, with human review built in at every step.

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed. Company decisions are never recorded automatically.

Getting Started

How the cockpit works

LandonOS runs entirely in your browser. There is no separate login server and nothing to install — your work is saved locally and is there when you return. Everything is organized as a set of cockpit instruments, each covering one stage of responsible research.

HOW TO READ THIS GUIDE

- Each page that follows covers one instrument: what it is, what you can do with it, and the exact steps to use it.
- The cockpit display name is shown first, with the underlying module in the line beneath it.
- Prefer to watch instead of read? The app also offers a narrated walkthrough and a guided tour from the help menu.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed.
Company decisions are never recorded automatically.

Performance Cockpit

Command Center

The Performance Cockpit is your daily dashboard and the first screen you see. A cluster of circular gauges reads out the health of your research at a glance, with live tiles below for everything currently in motion. Treat it as your pre-drive instrument check: glance here first, then dive into whatever needs attention.

WHAT YOU CAN DO

- Instrument gauges for source quality, report readiness, research velocity, and human-review risk.
- Live tiles for open requests, blocked items, reports in progress, and completed handoffs.
- Today's priorities, review alerts, and quick actions that launch the most common tasks.

HOW TO USE IT

- 1 Open the Cockpit from the top of the sidebar — it loads automatically when you start a session.
- 2 Read the four gauges first: their arcs sweep from healthy to at-risk so you can spot trouble in one look.
- 3 Scan the live tiles for anything blocked or awaiting review.
- 4 Use a quick action to jump straight into the task that needs you most.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed.
Company decisions are never recorded automatically.

Research Engine

Guided Research Builder

Every project fires up here. Scoping the question before involving any AI is what keeps the rest of the workflow disciplined — a clear brief now prevents wasted laps later. The Engine walks you through framing the decision, the evidence, and the boundaries before a single source is gathered.

WHAT YOU CAN DO

- Define the precise decision the research is meant to support.
- List the sources the work will require up front.
- Record what must never be assumed, and flag requests that need human review.

HOW TO USE IT

- 1 Start a new project and write, in one sentence, the decision it must support.
- 2 List the sources you expect to need before you involve any AI assistance.
- 3 Note any assumptions that are off-limits and mark items that require human review.
- 4 Save the brief to send it into the Track Map workflow.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed.
Company decisions are never recorded automatically.

Track Map

Research GPS

The Track Map guides you through the full ten-step workflow, from understanding the question all the way to a clean handoff, with a progress gauge to keep you on the racing line. It is the single source of truth for where each mission stands, so nothing is skipped and nothing stalls silently.

WHAT YOU CAN DO

- Pick an active mission and drive through each step in order.
- Mark each step Complete, In Progress, or Needs Help.
- Your progress keeps the cockpit gauges and counts accurate in real time.

HOW TO USE IT

- 1 Select an active mission to open its ten-step route.
- 2 Work each step in order so context carries forward cleanly.
- 3 Set every step to In Progress, Complete, or Needs Help as you go.
- 4 Watch the progress gauge — it feeds the cockpit counts and keeps your status honest.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed.
Company decisions are never recorded automatically.

Source Garage

Source Vault

The Source Garage is where evidence is stored and rated, so trust is never assumed. Each source carries a clear label, and the Garage actively warns you when a conclusion is resting on unverified ground. This is the backbone of the compliance guardrail — verified evidence in, defensible reports out.

WHAT YOU CAN DO

- Official and internal records are treated as trusted.
- AI drafts and unknown sources are clearly flagged as unverified.
- A risk banner warns when a request relies only on unverified sources.

HOW TO USE IT

- 1 Add each source and tag it Official, AI Draft, or Unknown.
- 2 Confirm an unknown or AI-drafted source against an official record before you rely on it.
- 3 Promote a source to trusted only once that check is done.
- 4 Clear the risk banner by replacing unverified sources before the report goes out.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed. Company decisions are never recorded automatically.

Brief Builder

Report Builder

The Brief Builder assembles a leadership-ready report and scores how ready it is to share. Rather than guessing whether a report is complete, you get an objective readiness score that flags exactly what is still missing — so the work that leaves your hands is already defensible.

WHAT YOU CAN DO

- Capture the objective, executive summary, findings, risks, and recommendation.
- Attach the sources that back each finding.
- The readiness score warns you when official sources or answers are missing.

HOW TO USE IT

- 1 Fill in the objective, executive summary, findings, risks, and recommendation.
- 2 Attach a backing source to every finding so each claim is traceable.
- 3 Read the readiness score and open each warning it raises.
- 4 Resolve every warning before you move the report to handoff.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed.
Company decisions are never recorded automatically.

Company Brain Sync

RoseOS

RoseOS is the company brain and your research co-driver. Ask it what is on record, and propose updates that stay drafts until a human approves them. It keeps institutional knowledge in one place while making sure nothing changes the official record without a person signing off.

WHAT YOU CAN DO

- Query the decision log, requirements registry, and automation registry.
- Draft proposed updates to core records.
- Nothing is recorded automatically — every change requires human review.

HOW TO USE IT

- 1 Ask RoseOS what is already on record before assuming anything.
- 2 Use its answers to fill gaps in your research instead of re-discovering known facts.
- 3 Draft a proposed update when you find something the brain should know.
- 4 Send the draft for human review — it is recorded only after a person approves it.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed.
Company decisions are never recorded automatically.

Garage Rewards & Driver Training

Growth & Rewards

As you work, Garage Rewards turns real research skill into points and badges, and Driver Training builds the skills behind them — both tracked on their own gauges. Progress reflects genuine milestones, not busywork, so the rewards always map back to better, safer research.

WHAT YOU CAN DO

- Points and badges reflect genuine research milestones, from Source Finder upward.
- Driver Training builds skills in using AI responsibly and verifying sources.
- A level gauge and roadmap show how close you are to the next milestone.

HOW TO USE IT

- 1 Complete real research milestones to earn points and badges.
- 2 Take Driver Training lessons to sharpen verification and responsible-AI habits.
- 3 Check the level gauge and roadmap to see exactly what unlocks the next level.
- 4 Revisit a lesson any time — reopening one simply removes its points so progress stays honest.

HUMAN REVIEW REQUIRED

Compliance guardrail: AI output is a draft only until it is source-checked and human-reviewed. Company decisions are never recorded automatically.